



CS 535 Deep Learning Assignment 2

**Mohamad Hosein Danesh
ID: 933619683**

February 13, 2019

1) The function that evaluates the trained model is implemented in the *MLP* class, defined as below:

```
def evaluate(self, x, y, batch_size, l2_penalty)
```

In this function, *x* is the input to the network. *y* is the output of the network. *batch_size* is the batch size of data given to the network at once. *l2_penalty* is the l2 regularization term.

The functions that compute the subgradients using backpropagation are implemented in *SigmoidCrossEntropy*, *ReLU*, and *LinearTransformation* classes. Each backward function has its own definition with respect to the class they are in. Below, every one of them will be shown:

```
class SigmoidCrossEntropy(object):  
def backward(self, y, grad_output)
```

```
class ReLU(object):  
def backward(self, grad_output)
```

```
class LinearTransform(object):  
def backward(self, grad_output)
```

2) In the function *train()* which is defined in the *MLP* class, the weight and biases will be updated after each iteration of mini-batches. Also, momentum has been used during weights and biases update.

3) The best performance I got was as following:

[EPOCH 100/100] Train Accuracy: 96.228 Train Loss: 0.375 Validation Accuracy: 81.232
Validation Loss: 2.267
Test Accuracy: 82.448 Test Loss: 2.128

To achieve this, I used the following hyperparameters:

hidden_units: 512
num_epochs: 100
learning_rate: 0.001
batch_size: 128
momentum: 0.6


4) With respect to the hyperparameters in part 3, the requested values are printed as follows(outputs from epoch 11 to 90 are omitted):

[EPOCH 1/100]	Train Accuracy: 67.598 Train Loss: 4.526	Validation Accuracy: 75.625 Validation Loss: 3.036
[EPOCH 2/100]	Train Accuracy: 70.333 Train Loss: 3.98	Validation Accuracy: 74.401 Validation Loss: 3.196
[EPOCH 3/100]	Train Accuracy: 72.102 Train Loss: 3.604	Validation Accuracy: 75.556 Validation Loss: 2.9
[EPOCH 4/100]	Train Accuracy: 74.639 Train Loss: 3.129	Validation Accuracy: 76.802 Validation Loss: 2.767
[EPOCH 5/100]	Train Accuracy: 75.796 Train Loss: 2.935	Validation Accuracy: 77.24 Validation Loss: 2.679
[EPOCH 6/100]	Train Accuracy: 76.787 Train Loss: 2.76	Validation Accuracy: 77.418 Validation Loss: 2.644
[EPOCH 7/100]	Train Accuracy: 77.587 Train Loss: 2.631	Validation Accuracy: 76.855 Validation Loss: 2.719
[EPOCH 8/100]	Train Accuracy: 78.312 Train Loss: 2.516	Validation Accuracy: 75.874 Validation Loss: 2.886
[EPOCH 9/100]	Train Accuracy: 78.933 Train Loss: 2.419	Validation Accuracy: 76.151 Validation Loss: 2.853
[EPOCH 10/100]	Train Accuracy: 79.598 Train Loss: 2.322	Validation Accuracy: 76.444 Validation Loss: 2.825
.		
. (omitted outputs)		
.		
[EPOCH 90/100]	Train Accuracy: 95.854 Train Loss: 0.412	Validation Accuracy: 81.096 Validation Loss: 2.282
[EPOCH 91/100]	Train Accuracy: 95.9 Train Loss: 0.408	Validation Accuracy: 81.112 Validation Loss: 2.28
[EPOCH 92/100]	Train Accuracy: 95.944 Train Loss: 0.403	Validation Accuracy: 81.125 Validation Loss: 2.279
[EPOCH 93/100]	Train Accuracy: 95.987 Train Loss: 0.399	Validation Accuracy: 81.14 Validation Loss: 2.277
[EPOCH 94/100]	Train Accuracy: 96.029 Train Loss: 0.395	Validation Accuracy: 81.157 Validation Loss: 2.275
[EPOCH 95/100]	Train Accuracy: 96.07 Train Loss: 0.391	Validation Accuracy: 81.173 Validation Loss: 2.274
[EPOCH 96/100]	Train Accuracy: 96.111 Train Loss: 0.387	Validation Accuracy: 81.186 Validation Loss: 2.272
[EPOCH 97/100]	Train Accuracy: 96.151 Train Loss: 0.383	Validation Accuracy: 81.2 Validation Loss: 2.27

[EPOCH 98/100] Train Accuracy: 96.189 Train Loss: 0.379 Validation Accuracy: 81.215
Validation Loss: 2.269
[EPOCH 99/100] Train Accuracy: 96.228 Train Loss: 0.375 Validation Accuracy: 81.232
Validation Loss: 2.267
Test Accuracy: 82.448 Test Loss: 2.128


Also there are diagrams of the output values(all the outputs are shown in diagrams):

Results(hidden_units: 512, num_epochs: 100, learning_rate: 0.001, batch_size: 128, momentum: 0.6)




5) Test accuracy for trying different batch sizes:


Test Accuracy(hidden_units = 512, learning_rate = 0.001, num_epochs = 10)



Test Accuracy(hidden_units = 512, learning_rate = 0.001, num_epochs = 10)




Test Accuracy(hidden_units = 512, learning_rate = 0.001, num_epochs = 10)




And by putting these three diagram into one we will have the following diagram:

Test Accuracy(hidden_units = 512, learning_rate = 0.001, num_epochs = 10)




Test accuracy for trying different learning rates:


Test Accuracy(batch_size = 32, hidden_units = 512, num_epochs = 10)



Test Accuracy(batch_size = 32, hidden_units = 512, num_epochs = 10)




Test Accuracy(batch_size = 32, hidden_units = 512, num_epochs = 10)




And by putting these three diagram into one we will have the following diagram:

Test Accuracy(batch_size = 32, hidden_units = 512, num_epochs = 10)




Test accuracy for trying different hidden units:


Test Accuracy(batch_size = 32, learning_rate = 0.005, num_epochs = 10)



Test Accuracy(batch_size = 32, learning_rate = 0.005, num_epochs = 10)

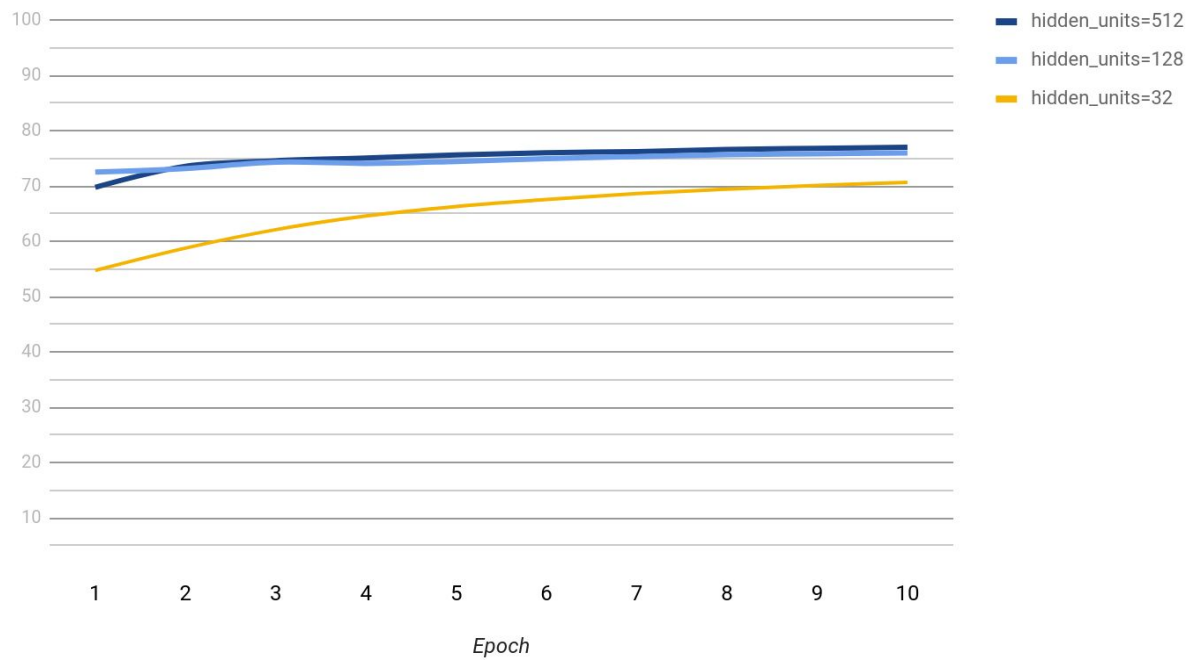


Test Accuracy(batch_size = 32, learning_rate = 0.005, num_epochs = 10)



And by putting these three diagram into one we will have the following diagram:

Test Accuracy(batch_size = 32, learning_rate = 0.005, num_epochs = 10)



6) Although it is not a convolutional neural network and has only one hidden layer, it performs reasonably well. At final epochs, it gets around 96% accuracy in the training set. But it only gets around 82% accuracy in the test set. In my opinion that is because the model is overfitted over the training data.

According to the experiments done in the part 5 of the assignment, different batch size does not have much impact on the test accuracy, eventually. Nevertheless, the batch size equal to 128 got the most test accuracy among different values. They all got closed to the 75% accuracy.

Different learning rates also have little impact on the final accuracy. As it is obvious from the diagrams of part 5, the three different settings all got to about 76% accuracy. Although learning rate equals to 0.005 got the most accuracy among others. Regarding the number of hidden units, according to experiments run on part 5, it has the most impact on the model's accuracy. The bigger the hidden layer, the more accurate it gets. It got to around 77% accuracy with 512 hidden neurons.